

Stage-1 Pre-Filter Screening Prompt

Replication artifact for "LLM-Based Test Oracles: Source-of-Authority Taxonomy"

Ali Hassaan Mughal and Muhammad Bilal

LLM PRE-FILTER SCREENING PROMPT (Stage 1)

Systematic review: "LLM-Based Test Oracles: A Source-of-Authority Taxonomy"

Model: Claude Opus 4.8 | Decoding: default settings, no sampling parameters overridden | Run: June 2026

Criteria below are the pre-registered eligibility criteria exactly as fixed in the review protocol and stated in the manuscript (Section IV, Eligibility criteria).

You are screening bibliographic records for a systematic literature review titled "LLM-Based Test Oracles: A Source-of-Authority Taxonomy."

Review scope: primary studies in which a large language model (LLM) produces, or itself acts as, a test oracle -- the component of a test that decides whether observed software behavior is correct.

You are given ONLY the record's TITLE and ABSTRACT. Judge the record against the pre-registered eligibility criteria below and return a single decision. Do not use any knowledge beyond the title and abstract provided.

INCLUSION CRITERIA (a record is included only if ALL hold):

- I1 (Oracle focus): The work concerns a test oracle, that is, the mechanism that decides expected behavior or pass/fail (assertions, expected outputs, properties or invariants, metamorphic relations, differential comparison, or human judgment), rather than only test-input generation or coverage.
- I2 (LLM involvement): An LLM generates, infers, or determines the oracle, or the study empirically examines how the oracle is determined or behaves in an LLM-based testing pipeline.
- I3 (Window): Published between January 2021 and May 2026.
- I4 (Form): At least four pages, written in English, with retrievable full text, peer reviewed and indexed (conference, journal, or early access).

EXCLUSION CRITERIA (a record is excluded if ANY holds; give the reason code):

- E1: Testing OF LLMs or natural language processing (NLP) systems (the LLM is the system under test, for example bias, hallucination, jailbreak, or robustness testing), rather than LLMs used FOR testing software.
- E2: LLM test or code generation that does not foreground an oracle (input-only, coverage-only, prefix-only, or program-repair-only).
- E3: Hardware-verification assertion or oracle generation, for example in SystemVerilog, register-transfer level (RTL), or very large-scale integration (VLSI).
- E4: Secondary studies (surveys, systematic reviews, mapping studies).
- E5: Non-peer-reviewed preprints.
- E6: Editorials, keynotes, posters, tool demonstrations, tutorials, extended abstracts, items under four pages, non-English items, and non-retrievable items.

DECISION RULE:

- Return exactly one decision: INCLUDE, MAYBE, or EXCLUDE.
- If the title and abstract are ambiguous or insufficient to decide confidently, return MAYBE and retain the record for full-text assessment. Never EXCLUDE on uncertainty.
- If EXCLUDE, give the single most specific applicable reason code (E1-E6).

OUTPUT FORMAT (one record):

decision: <INCLUDE | MAYBE | EXCLUDE>
reason_code: <E1-E6, or none>
rationale: <one line>
